

Efficient Optimal Search under Expensive Edge Cost Computation

Masataro Asai, Akihiro Kishimoto, Adi Botea, Radu Marinescu, Elizabeth M. Daly and Spyros Kotoulas
The University of Tokyo, IBM Research Ireland


Motivation: Domains where edge-costs are not known apriori

Example 1: TSP w/ unknown edge costs

- $O(N^2)$ edge cost computations
- N : number of cities
- TSP itself is NP-complete

Example 2: Multiple Worker Routing Problem (MWRP)

A problem of finding the journey plans for workers (doctors, nurses) to attend the patients across the city, using multi-modal journey (using buses, trains, walking, etc...)




- $O(N^2T)$ edge cost computations
- N : number of locations
- T : number of possible start time (e.g. 1 day = $24 \times 60 \times 60$ sec)
- Optimally solving MWRP is NP-hard (by reduction from SMTTP [Du and Leung, 1990])

city	nodes	segments	stops	routes	trips/day
Dublin	301,638	319,846	4,739	120	7,308
Montpellier	152,949	161,768	1,297	36	3,988
Rome	522,529	566,400	8,896	391	39,422

Table 1: Statistics of transport data used in experiments.

Trivial approach with A*: Compute the edge costs on demand

- Compute the edge costs during the search
- Evaluate the edge cost using an external solver DIJA [Botea et al., 2013] when computing the $g()$ value of A^*
- The result of solving DIJA is cached (the same edge is never computed twice)

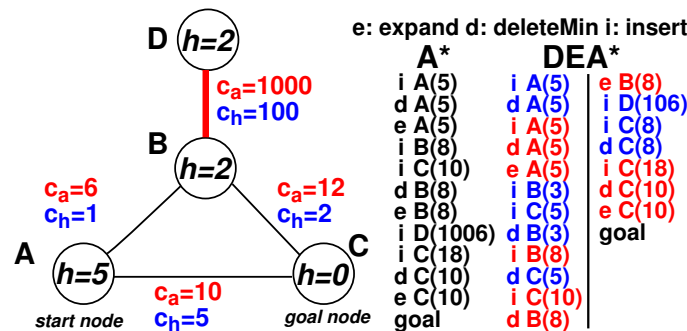


- Even with caching, most of the runtime is spent on calling DIJA solver

Delayed Expansion A^* (DEA*):

- The **actual edge cost c_a** is too expensive to compute. In many cases, c_a is unnecessary.
- Delay the evaluation of c_a by using the **lower-bound of the edge cost, c_h** , provided by the external solver
- Insert the same node twice, f -value computed once by c_h , then by c_a

Pathological case for A^* : DEA* can avoid evaluating the edge **BD**



Evaluating DEA* on MWRP

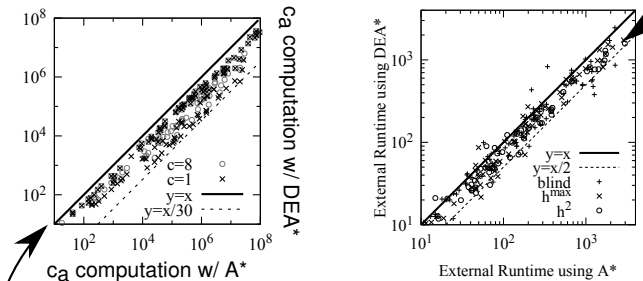
A^*/DEA^* using Blind / h^{\max} / h^2 heuristics

DEA* with $h^1 h^2$: use h^{\max} in the first expansion, use h^2 in the second expansion

worker		Blind		h^{\max}		h^2		$h^1 h^2$
City	W	A*	DEA*	A*	DEA*	A*	DEA*	DEA*
Total		50	59	68	77	81	83	85
Dublin	12	14	15	15	21	19	22	22
	6	2	4	8	8	12	12	11
Montpellier	12	16	15	17	17	17	15	18
	6	1	4	3	5	3	4	5
Rome	12	15	18	18	19	19	19	20
	6	2	3	7	7	11	11	9

Number of solved instances

- Improved overall performance
- In some cases **x2 OPEN list operation** can be a bottleneck, losing to A^*
- Time spent for calling DIJA has been reduced
- Results in the smaller overall runtime



DEA* on Classical Planning

- Evaluate the domain-independent performance
- Standard IPC benchmark assumes the edge costs are given, so we "simulate" the case where c_a is expensive & calculated dynamically
- Simulated DEA*: $c_h = c_a - C$, C : constant value
- DEA* reduced the c_a computation (up to x30)