

Classical Planning in Deep Latent Space: From Unlabeled Images to PDDL (and back)

Masataro Asai,
Alex Fukunaga
The University of Tokyo

1. Background

- Symbols: Labels for distinguishing entities
- Symbols in PDDL: Propositions (handempty)
Objects A, B, C, D
Predicates (clear ?x)
Actions (pick-up ?x)
- Existing Action Model Learning (AML) systems ALL require symbolic (or near-symbolic) inputs
- Cannot directly handle complex unstructured high-dimensional data e.g. images, audio

Action-Relation Modelling System [Yang et al 07]

input: symbolic relations (predicates, actions)

Semi-MDP → Classical Planning [Konidaris et al 14, 15]

Not entirely subsymbolic: converts probabilistic → propositional model

Symbolic action labels: given (SMDP option e.g. move, interact)

Input sensors: given, structured, low-dim

(set of 33 real/int variables with distinct meanings: x/y distance, light level, whether monkey cries)

Learning from Observation: [Argall et al 09, Mourao et al 12]

Noisy, incomplete, but symbolic states/actions

Learns the preconditions from state/action sequence

Object/predicate/action labels: given

Learning from Video for board game:

Images with strong assumptions (almost symbol)


e.g. 3x3 Ellipse Detector [Barbu et al 10; Kaiser12; Kirk&Laird16]

almost immediately provides propositions

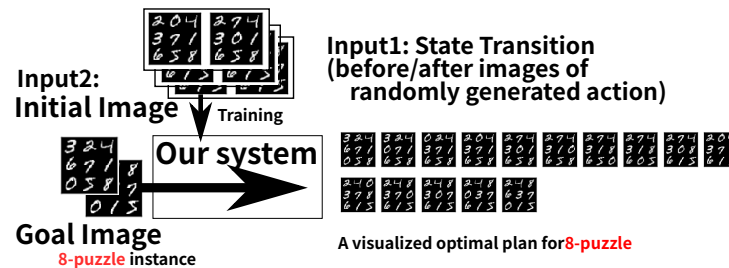
2. Our Objectives

Visually presented classical planning domains

The system should have no idea that this is an 8-puzzle



- with NO prior assumptions/symbolic descriptions ("grids", "tiles", "disks", "move", "toggle") i.e. domain-independent
- with NO expert plan traces



Existing AML systems cannot handle raw images

- High-dimensional input (42x42 pixels = 1764)
- Each pixel does not have a significant meaning
- "Meaning" emerges from nonlin. entanglements
- No simple detector (e.g. stone existence)
- Need Robustness & Generalization for noisy inputs

3. OPEN PROBLEM: Convert images/ arbitrary unstructured data to/from symbols

Generate the symbolic inputs for AML Systems!


Raw data → Symbolic data

→ Existing AML methods

→ Symbolic Planning → Symbolic plan → Raw output

4. Solution: State AutoEncoder

- A neural network that learns a bidirectional mapping $z = \text{Encode}(x)$ and $x = \text{Decode}(z)$ between a subsymbolic input x and a discrete boolean vector z
- The only specification is the number of bits: An upper bound of state encoding length $\log_2(|S|)$



- For training, optimize the loss via Stoc. Grad. Descent
Minimize reconstruction loss $\|x - \hat{x}\|$ (Binary crossentropy etc.)
+ variational loss $KL(Z, \text{Categorical}(49))$
(KL divergence: distance between distributions)

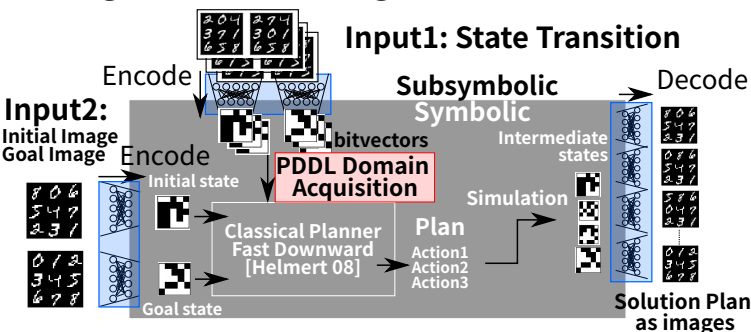
SAE: A Variational AutoEncoder w/ Gumbel-Softmax

- Allows NNs to learn a categorical/discrete distribution
 $z_i = \text{Softmax}(g_i + \log \pi_i / \tau)$
 π_i : Probability for Category i τ : Annealing Temperature
 g_i : A random sample from $\text{Gumbel}(0, 1)$ [Gumbel et.al, 1954]

Number of categories = 2 : Propositional variables

- the "meaning" of the propositions may or may not correspond to human intuitions about the domain

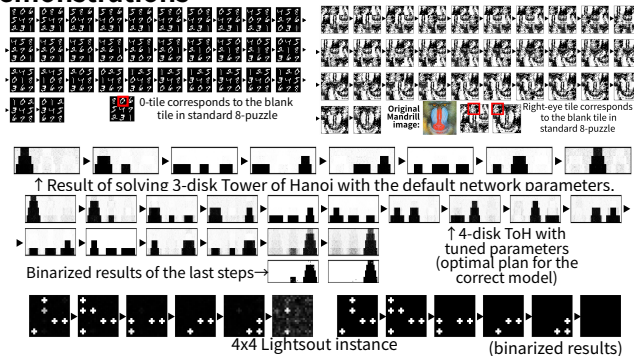
5. Image-based planning architecture: LatPlan




A First Implementation: LatPlan α

- SAE: Trained w/ smaller subset of states (cross validated, SAE learns a generalized/robust state encoding)
- Trivial AML method: action = individual state transition we provide all state transitions (i.e. no action generalization)
- Future work: adapt existing AML method

Demonstrations Optimal solutions to 8-puzzle instances



Handling of Noisy Inputs (benefit of using NN)



6. Domain-Independent Search Heuristics are Effective in Latent-Space Planning

Somewhat surprising result: Heuristics reduce search in domains based on features generated by neural networks

- Exploit common characteristics (e.g. abstraction)
- PDBs can underperform blind search in some standard IPC domains (Edelkamp12)

Heuristics	Search Time (sec)	Expansion
MNIST 8-puzzle (6 different instances)		
blind	1.904	193924
pdb	1.780	109096
blind	1.751	201156
pdb	1.508	111642
blind	1.657	186767
pdb	1.215	84561
blind	1.514	183336
pdb	1.474	82518
blind	1.460	169907
pdb	0.685	52084
blind	1.489	130863
pdb	0.382	26967
Hanoi (4 peg)		
blind	0.0008	55
pdb	0.0006	17
LightsOut (4x4)		
blind	0.0159	952
pdb	0.0013	27
Spiral LightsOut (3x3)		
blind	0.0040	522
pdb	0.0026	214
Mandrill 8-puzzle		
blind	2.759	335378
pdb	1.113	88851

Example latent-space domain

Example with $|z| = 25$ bits

```
(define (domain latent)
  (:requirements :strips :negative-preconditions)
  (:predicates (z0) (z1) (z2) (z3) (z4) ... (z24))
  (-action at00001001011011100011110000100010111111001111
    :parameters () :precondition
    (and (z0) (not (z1)) (not (z2)) (not (z3)) (not (z4)) (not (z5)) (z6) (not (z7))
      (not (z8)) (z9) (not (z10)) (z11) (z12) (not (z13)) (z14) (z15) (z16) (z17)
      (not (z18)) (not (z19)) (not (z20)) (z21) (z22) (z23) (z24))
    :effect (and (z5) (not (z6)) (z13) (z20)) ...)
```

pre: before-state
eff: state diff